

# Data Frame Manipulation

Juliana

October 31, 2017

## Introduction

The purpose of this post is to expand upon the topic of data frame manipulation. I go about doing this by focusing on particular sets of functions. The topic of data frame manipulation is itself a very broad topic that encompasses an innumerable amount of examples and specific applications; I, therefore, chose a small variety of functions that I felt are important for any R user to know.

## Discussion

I would first like to start this post by discussing a data frame application that I came across while re-looking into the base R function `aggregate()`.

When learning R, I think it is important to be aware of both the differences and similarities between base R functions and functions that are provided by R packages.

This first example I will work with demonstrates a pretty simple application of the `aggregate()` function. Therefore, I will be reviewing this particular example not to show how `aggregate()` works, per se, but to demonstrate how `group_by()` and `summarise()` can be used to produce the same data frame. I think this is important to point out simply because it's easy to lose sight of the fact—especially if one always uses “dplyr” functions—that `aggregate()` can essentially do the same thing that `group_by()` and `summarise()` can do together. The `aggregate()` function has not been used too much in Stat 133, so I wanted to make this particular relationship between a base R function and “dplyr” functions very clear.

*Relationship Between base R Function `aggregate()` and “dplyr” Functions `group_by()` and `summarise()`*

Let's start by creating a simple data frame. Don't forget the `stringAsFactors = FALSE`!

```
my_pets <- data.frame(pets = c(rep("dog", 4),
                              rep("cat", 6),
                              rep("bunny", 10)),
                     stringsAsFactors = FALSE)
```

my\_pets

```
##      pets
## 1     dog
## 2     dog
## 3     dog
## 4     dog
## 5     cat
## 6     cat
## 7     cat
## 8     cat
## 9     cat
## 10    cat
## 11  bunny
## 12  bunny
## 13  bunny
## 14  bunny
## 15  bunny
## 16  bunny
## 17  bunny
## 18  bunny
## 19  bunny
## 20  bunny
```

Now that we have created the data frame, let's suppose that we want to separate our pets into groups and count how many of each pet we have. This is where the `aggregate()` function will come in.

```
num_of_pets <- aggregate(my_pets, by = list(unique_pets = my_pets$pets), FUN = length)
```

num\_of\_pets

```
##      unique_pets pets
## 1         bunny  10
## 2          cat    6
## 3          dog    4
```

Voila! We get a new and improved data frame in which we can clearly see that we have 10 bunnies, 6 cats, and 4 dogs.

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##      filter, lag
```

```
## The following objects are masked from 'package:base':
##
## intersect, setdiff, setequal, union
```

```
num_of_pets2 <- summarise(group_by(my_pets, pets), num_of_each = length(pets))

num_of_pets2
```

```
## # A tibble: 3 x 2
##   pets num_of_each
##   <chr>      <int>
## 1 bunny         10
## 2 cat           6
## 3 dog           4
```

Great! We see that we have gotten the same output. To summarise, we can think of the relationship between these functions as such: the “by” parameter in the `arrange()` function is equivalent to the `group_by()` function in “dplyr” while the “FUN” parameter in `arrange()` is equivalent to `summarise()`.

### Creating Dates and Extending our Understanding of `aggregate()`

I will now be moving on to another example involving a data frame containing dates and which also involves an application of the `aggregate()` function. Before doing so, I would like to give credit for this example to David Kun, as it was the example he wrote in his own article on `aggregate()` that inspired the example presented below. A link to his article can be found in the References section of this blog post.

As you will see below, the function `as.Date()` can be used to generate dates.

```
# Generating a data frame with 1 year's worth of dates

my_dates <- data.frame(dates = as.Date("2017-01-12", format = "%Y-%m-%d") + 0:365)

head(my_dates, 10)
```

```
##           dates
## 1 2017-01-12
## 2 2017-01-13
## 3 2017-01-14
## 4 2017-01-15
## 5 2017-01-16
## 6 2017-01-17
## 7 2017-01-18
## 8 2017-01-19
## 9 2017-01-20
## 10 2017-01-21
```

I find this example interesting for two reasons. One, because it shows how it is actually quite simple to generate a data frame full of dates. And two, because it shows us that we can generate a year’s worth of dates by simply adding a “+” after `as.Date()` and specifying the number of days that we would like to have in our data frame. In this case, I chose exactly 365.

Now, let’s use `aggregate()` on `my_dates`.

```
min_date <- aggregate(my_dates, by = list(month = substr(my_dates$dates, 1, 7)), FUN = min)

min_date
```

```
##      month      dates
## 1 2017-01 2017-01-12
## 2 2017-02 2017-02-01
## 3 2017-03 2017-03-01
## 4 2017-04 2017-04-01
## 5 2017-05 2017-05-01
## 6 2017-06 2017-06-01
## 7 2017-07 2017-07-01
## 8 2017-08 2017-08-01
## 9 2017-09 2017-09-01
## 10 2017-10 2017-10-01
## 11 2017-11 2017-11-01
## 12 2017-12 2017-12-01
## 13 2018-01 2018-01-01
```

From the data frame displayed above, we can see that we have grouped the data frame by the month of each individual year and next to each month displayed the earliest possible date for that month. Notice that the date next to “2017-01” and “2018-01” are different. This is because our starting date was “2017-01-12,” meaning that the earliest date in `my_dates` for the month of January in 2017 is “2017-01-12”, whereas the earliest date in January of 2018 is “2018-01-01”. In order to group by month and year, we use the `substr()` function. The 1 and 7 that are written as `substr()` parameters tell R to include characters 1 through 7 of each date, which encompass both the year and month. The `substr()` function in this case is very helpful. Without it, we wouldn’t be able to group by month since each date when taken as a whole is unique.

Let’s try to generate the same data frame using `summarise()` and `group_by()` like we did above!

```
# Generating a data frame with the earliest date for each month

summarise(group_by(my_dates, month = substr(dates, 1, 7)), dates = min(dates))
```

```
## # A tibble: 13 x 2
##   month      dates
##   <chr>    <date>
## 1 2017-01 2017-01-12
## 2 2017-02 2017-02-01
## 3 2017-03 2017-03-01
## 4 2017-04 2017-04-01
## 5 2017-05 2017-05-01
## 6 2017-06 2017-06-01
## 7 2017-07 2017-07-01
## 8 2017-08 2017-08-01
## 9 2017-09 2017-09-01
## 10 2017-10 2017-10-01
## 11 2017-11 2017-11-01
## 12 2017-12 2017-12-01
## 13 2018-01 2018-01-01
```

Awesome! We see once again how `aggregate()` and the functions `group_by()` and `summarise()` are really similar and allow us to produce the same data frame!

Ok, now for another `aggregate()` example.

Let's create two data frames. The first will represent the category into which a scoring of a particular vacation spot falls. The second data frame will represent the actual score given to a vacation spot by someone who went there.

```
location_ratings <- data.frame(rankings = c(rep("breathtaking", 7), rep("so-so", 5), rep("ugly", 12)), stringsAsFactors = FALSE)

raw_scores <- data.frame(scores = c(rnorm(7, mean = 9.8, sd = 0.3), rnorm(5, mean = 5, sd = 0.4), rnorm(12, mean = 1.5, sd = 0.5)), stringsAsFactors = FALSE)
```

Let's create a data frame in which the scores are grouped into categories.

```
categorized_scores <- aggregate(raw_scores, by = location_ratings, FUN = function(raw_scores){
  max(raw_scores) * 2
})

categorized_scores
```

```
##   rankings      scores
## 1 breathtaking 20.614320
## 2      so-so 11.199507
## 3      ugly  6.075541
```

The main point of this example is to show that you **don't** have to use the same data frame for the `aggregate()` function and its "by" parameter. In other words, the column variable or variables by which you decide to group does not have to come from the data frame that is used as the first argument of the `aggregate()` function. As shown above, you can also create your own function for the FUN parameter.

### Untidy Data and "tidyr" Functions

I will now be moving on to talk about what, in the R community, is referred to as untidy data and demonstrate the application of a couple of functions from the "tidyr" package created by Hadley Wickham.

Because in the real world the data files that one works with will not always be in a tidy format in which rows serve as observations and columns as variables, it is important to learn how to transform data into a tidy format.

Let's start by loading the packages we'll be working with and importing the data contained within `treats.csv`. In the data we'll be looking at, there will be 4 individuals, the different types of treats they consume (cake, chocolate, and ice cream), the number of servings they have of each (1 or 2), and the number of miles they run when they consume 1 or 2 servings. In order to display the data in a tidy format, we'd ideally like to have one column for each of the following: names, treats, servings, and miles run.

```
library(tidyr)
```

```
## Warning: package 'tidyr' was built under R version 3.4.2
```

```
library(dplyr)
```

```
treats_data <- read.csv("../Post1/treats.csv", stringsAsFactors = FALSE)
treats_data
```

```
##   name cake_1 cake_2 choco_1 choco_2 icream_1 icream_2
## 1 Samara   3.0   6.00    1.0     4     5.0     9
## 2 Nancy    0.0   2.00    0.5     3     1.0     4
## 3 Sheldon  0.5   0.75    1.0     2     1.5     2
## 4 Brianna  4.0   6.00    2.0     5     6.0    10
```

After importing `treats.csv`, we can see that the data is not in a tidy format. To begin with, there isn't a miles column to indicate the number of

miles that each individual runs. Instead, every row is made up exclusively of values representing the number of miles run by each individual. We also see that the type of treat and number of servings have been combined in the header row. Additionally, instead of having just one column or variable for all of the treats, we see that each treat-serving combination serves as a column of its own.

Let's start tidying up the data!

This is where the `gather()` function comes in.

```
treats_data2 <- gather(treats_data, key = treat_serving, value = miles, cake_1:icream_2)
```

```
treats_data2
```

```
##      name treat_serving miles
## 1  Samara      cake_1    3.00
## 2   Nancy      cake_1    0.00
## 3 Sheldon      cake_1    0.50
## 4 Brianna      cake_1    4.00
## 5  Samara      cake_2    6.00
## 6   Nancy      cake_2    2.00
## 7 Sheldon      cake_2    0.75
## 8 Brianna      cake_2    6.00
## 9  Samara      choco_1    1.00
## 10 Nancy      choco_1    0.50
## 11 Sheldon      choco_1    1.00
## 12 Brianna      choco_1    2.00
## 13 Samara      choco_2    4.00
## 14 Nancy      choco_2    3.00
## 15 Sheldon      choco_2    2.00
## 16 Brianna      choco_2    5.00
## 17 Samara      icream_1    5.00
## 18 Nancy      icream_1    1.00
## 19 Sheldon      icream_1    1.50
## 20 Brianna      icream_1    6.00
## 21 Samara      icream_2    9.00
## 22 Nancy      icream_2    4.00
## 23 Sheldon      icream_2    2.00
## 24 Brianna      icream_2   10.00
```

As you can see, the `gather()` function produces what, in the R help documentation, are referred to as “key-value pairs”. What we set equal to key and value become the names of the new columns we create, as we can see in the data frame above. The code that comes after the value assignment (in this case, “`cake_1:icream_2`”) will become the input to the new column created by the key assignment. In other words, what were previously column/variable names will become the values of one column, namely the column whose name has been assigned to key. This is why we see a column named “`treat_serving`” that contains all of the treat-serving combinations. The values that fall under the value column, which in our case is called “`miles`,” are those values that had initially been listed under the variables that now makeup the key column.

The data frame definitely looks better, but we still need to separate the type of treat from the number of servings consumed in a day. In order to do this, we will be using the `separate()` function.

```
treats_data3 <- separate(treats_data2, treat_serving, into = c("treat", "serving"), sep = "_")
```

```
treats_data3
```

```
##      name treat serving miles
## 1  Samara  cake       1    3.00
## 2   Nancy  cake       1    0.00
## 3 Sheldon  cake       1    0.50
## 4 Brianna  cake       1    4.00
## 5  Samara  cake       2    6.00
## 6   Nancy  cake       2    2.00
## 7 Sheldon  cake       2    0.75
## 8 Brianna  cake       2    6.00
## 9  Samara  choco      1    1.00
## 10 Nancy  choco      1    0.50
## 11 Sheldon  choco      1    1.00
## 12 Brianna  choco      1    2.00
## 13 Samara  choco      2    4.00
## 14 Nancy  choco      2    3.00
## 15 Sheldon  choco      2    2.00
## 16 Brianna  choco      2    5.00
## 17 Samara  icream     1    5.00
## 18 Nancy  icream     1    1.00
## 19 Sheldon  icream     1    1.50
## 20 Brianna  icream     1    6.00
## 21 Samara  icream     2    9.00
## 22 Nancy  icream     2    4.00
## 23 Sheldon  icream     2    2.00
## 24 Brianna  icream     2   10.00
```


We finally got the data frame that we wanted!

Now, that we've created a tidy data frame, we can easily generate a scatterplot using `ggplot2` that will show us the distribution of the miles run by each individual after having eaten a particular treat.

```
# Load the package ggplot2
library(ggplot2)

# Scatterplot showing the miles run by each individual

ggplot(treats_data3, aes(name, miles)) + geom_point(aes(col = treat)) + labs(x = "Names", y = "Miles")
```



### Data Frames: from Wide to Long and Long to Wide

In addition to switching from what could be called an untidy to a tidy format, we can also interpret what we just did as having gone from a “wide” to a “long” format. A data frame that has a wide format will have more than 1 measurement in each row. In the original data frame, `treats_data`, every row was made up of only measurements. In other words, every column was a measurement variable. With a long data frame, on the other hand, you have only one measurement in each row, whereas the rest of the columns represent categorical variables. Our last data frame, `treats_data3`, can be classified as a long data frame since the only measurement in each row is the number of miles run. Don’t get confused by the “serving” column. Even though this column is made up of numbers, it still qualifies as a categorical variable. In her article on data wrangling, Sharon Machlis provides a good way to think about it. She essentially says that if it doesn’t make sense to plot the values of a column on their own, then that column is more than likely not a measurement column. In the case of `treats_data3`, it wouldn’t make sense to plot the repeating 1’s and 2’s of the “serving” column; it would, however, make sense to plot the number of miles run, especially if the plots were split up into miles run for 1 serving and miles run for 2 servings.

Moving forward, I would like to demonstrate how to go from a “wide” data frame to a “long” one using the package “reshape2” and the function `melt()`.

```
# Let's start by importing the data we will use for our data frame

weight_loss <- read.csv("../Post1/weight_loss_data.csv", stringsAsFactors = FALSE)
```

```
## Warning in read.table(file = file, header = header, sep = sep, quote =
## quote, : incomplete final line found by readTableHeader on '../Post1/
## weight_loss_data.csv'
```

```
weight_loss
```

```
##      name start_year      gym starting_weight end_weight
## 1  Nancy      2015      UFC             180         155
## 2 Sheldon      2014  24_HrFit             220         185
## 3 Samara      2016 Planet_Fitness           170         135
## 4 Brianna      2013  24_HrFit             180         150
## goal_weight
## 1          150
## 2          180
## 3          135
## 4          140
```

From the data frame `weight_loss`, we can clearly see that we have more than 1 measurement column; in this case, there are 3. This is where the function `melt()` will come in! After installing the package “reshape2,” I will load it into my Rmd file.

```
# Load "reshape2" into Rmd file

library(reshape2)
```

```
## Warning: package 'reshape2' was built under R version 3.4.2
```

```
##
## Attaching package: 'reshape2'
```

```
## The following object is masked from 'package:tidyr':
##
## smiths
```

In the code below, I will demonstrate how to use the `melt()` function in order to get the desired “long” data frame.

```
weight_loss_long <- melt(weight_loss, id.vars = c("name", "start_year", "gym"), variable.name = "weight_info", value.name = "pounds")

weight_loss_long
```

```
##      name start_year      gym  weight_info pounds
## 1  Nancy      2015      UFC starting_weight  180
## 2 Sheldon    2014    24_HrFit starting_weight  220
## 3 Samara     2016 Planet_Fitness starting_weight  170
## 4 Brianna    2013    24_HrFit starting_weight  180
## 5 Nancy      2015      UFC      end_weight  155
## 6 Sheldon    2014    24_HrFit      end_weight  185
## 7 Samara     2016 Planet_Fitness      end_weight  135
## 8 Brianna    2013    24_HrFit      end_weight  150
## 9 Nancy      2015      UFC      goal_weight  150
## 10 Sheldon   2014    24_HrFit      goal_weight  180
## 11 Samara    2016 Planet_Fitness      goal_weight  135
## 12 Brianna   2013    24_HrFit      goal_weight  140
```

Awesome! We now have only one measurement column, namely “pounds”, as opposed to three! From the code above, we can see that we did not need to explicitly write out the column names that we intended to become the values of our new column “weight\_info”. The `melt()` function will automatically convert the column names that are not specified in the vector that is set equal to `id.vars` into the input for the new categorical column. The parameter `variable.name` allows us to name the new categorical column, while the parameter `value.name` allows us to do the same for the new measurement column.

```
# If you feel more comfortable explicitly writing out which column names you wish to become the input for the new categorical column "weight_info", you can use the following code, which includes the parameter measure.vars

weight_loss_long <- melt(weight_loss, id.vars = c("name", "start_year", "gym"), measure.vars = c("starting_weight", "end_weight", "goal_weight"), variable.name = "weight_info", value.name = "pounds")


weight_loss_long
```

```
##      name start_year      gym  weight_info pounds
## 1  Nancy      2015      UFC starting_weight  180
## 2 Sheldon    2014    24_HrFit starting_weight  220
## 3 Samara     2016 Planet_Fitness starting_weight  170
## 4 Brianna    2013    24_HrFit starting_weight  180
## 5 Nancy      2015      UFC      end_weight  155
## 6 Sheldon    2014    24_HrFit      end_weight  185
## 7 Samara     2016 Planet_Fitness      end_weight  135
## 8 Brianna    2013    24_HrFit      end_weight  150
## 9 Nancy      2015      UFC      goal_weight  150
## 10 Sheldon   2014    24_HrFit      goal_weight  180
## 11 Samara    2016 Planet_Fitness      goal_weight  135
## 12 Brianna   2013    24_HrFit      goal_weight  140
```

Now that our data frame is in a “long” format, let’s create another scatterplot like the one we created for `treats_data3`!

```
# Scatterplot showing the different weights of each individual

ggplot(weight_loss_long, aes(name, pounds)) + geom_point(aes(col = weight_info)) + labs(x = "Names", y = "Weight")
```




The “long” format of `weight_loss_long` allows us to generate a scatterplot that clearly shows how far each individual was from their target weight. It appears that Samara was the only one to hit her goal weight. We do not see a green dot for her, because her goal weight and end weight overlap.

We could also use `weight_loss_long` to produce a bar chart that compares the total number of pounds for `starting_weight`, `end_weight`, and `goal_weight`.

```
# Bar chart comparing the total in pounds for the different categories within weight_info

ggplot(weight_loss_long, aes(x = weight_info, y = pounds, fill = name)) + geom_bar(stat = 'identity')
```



Now let's try going from a long data frame to a wide one!

The function we'll be using this time is `dcast()`.

```
# This code will give us the graph we started with

weight_loss_wide <- dcast(weight_loss_long, name + start_year + gym ~ weight_info, value.var = "pounds")

weight_loss_wide
```

```
##      name start_year      gym starting_weight end_weight
## 1 Brianna      2013      24_HrFit          180        150
## 2 Nancy        2015         UFC          180        155
## 3 Samara       2016 Planet_Fitness          170        135
## 4 Sheldon     2014      24_HrFit          220        185
## goal_weight
## 1          140
## 2          150
## 3          135
## 4          180
```

The first argument of the function is the data frame you will be changing. The column names written before the ~ (and separated by a "+") are the columns you'd like to keep, while the column name placed after the ~ indicates the column whose unique values will be used to create new columns. The values for these new columns, in our case, will come from the column "pounds" (the measurement column), hence the assignment of this variable name to the parameter value.var.

Here one can see that what we did with the gather() function from the package "tidyr" was essentially the same as what we did with the melt() function from the package "reshape2"! The same goes for the functions spread() and dcast()! I did not provide an example for spread(), but it essentially does the same thing as dcast().

### Family of join() Functions

Now that we've talked about "tidyr" and "reshape2", I'd like to move on to discuss some of the joining functions provided by "dplyr". Merging has been very briefly touched upon in Stat 133, so I'd like to expand upon this.

I'd like to start off by saying that although I will be discussing the join functions provided by "dplyr", the same operations can be performed using the function merge(), which is a base R function.

I will be reviewing the following functions:

- inner\_join()
- semi\_join()
- left\_join()
- right\_join()
- full\_join()

Ok, let's start with inner\_join()!

In order to show how inner\_join() works, I will be using two data frames from above, weight\_loss\_long and treats\_data3. We will see that inner\_join() will return the rows that are in both weight\_loss\_long and treats\_data3 as well as all of the columns in each of those data frames. We know that the "name" column of each data frame contains the exact same names, so it is by the "name" variable that we are joining the two data frames.

```
head(inner_join(weight_loss_long, treats_data3), 26)
```

```
## Joining, by = "name"
```



```
##      name start_year      gym  weight_info pounds  treat
## 1   Nancy      2015      UFC starting_weight    180   cake
## 2   Nancy      2015      UFC starting_weight    180   cake
## 3   Nancy      2015      UFC starting_weight    180  choco
## 4   Nancy      2015      UFC starting_weight    180  choco
## 5   Nancy      2015      UFC starting_weight    180 icream
## 6   Nancy      2015      UFC starting_weight    180 icream
## 7  Sheldon      2014    24_HrFit starting_weight    220   cake
## 8  Sheldon      2014    24_HrFit starting_weight    220   cake
## 9  Sheldon      2014    24_HrFit starting_weight    220  choco
## 10 Sheldon      2014    24_HrFit starting_weight    220  choco
## 11 Sheldon      2014    24_HrFit starting_weight    220 icream
## 12 Sheldon      2014    24_HrFit starting_weight    220 icream
## 13 Samara      2016 Planet_Fitness starting_weight    170   cake
## 14 Samara      2016 Planet_Fitness starting_weight    170   cake
## 15 Samara      2016 Planet_Fitness starting_weight    170  choco
## 16 Samara      2016 Planet_Fitness starting_weight    170  choco
## 17 Samara      2016 Planet_Fitness starting_weight    170 icream
## 18 Samara      2016 Planet_Fitness starting_weight    170 icream
## 19 Brianna     2013    24_HrFit starting_weight    180   cake
## 20 Brianna     2013    24_HrFit starting_weight    180   cake
## 21 Brianna     2013    24_HrFit starting_weight    180  choco
## 22 Brianna     2013    24_HrFit starting_weight    180  choco
## 23 Brianna     2013    24_HrFit starting_weight    180 icream
## 24 Brianna     2013    24_HrFit starting_weight    180 icream
## 25 Nancy      2015      UFC      end_weight    155   cake
## 26 Nancy      2015      UFC      end_weight    155   cake
##      serving miles
## 1         1  0.00
## 2         2  2.00
## 3         1  0.50
## 4         2  3.00
## 5         1  1.00
## 6         2  4.00
## 7         1  0.50
## 8         2  0.75
## 9         1  1.00
## 10        2  2.00
## 11        1  1.50
## 12        2  2.00
## 13        1  3.00
## 14        2  6.00
## 15        1  1.00
## 16        2  4.00
## 17        1  5.00
## 18        2  9.00
## 19        1  4.00
## 20        2  6.00
## 21        1  2.00
## 22        2  5.00
## 23        1  6.00
## 24        2 10.00
## 25        1  0.00
## 26        2  2.00
```

In order to avoid getting the message that tells us we're joining by "name", we can include the "by" parameter in the inner\_join() function.

```
head(inner_join(weight_loss_long, treats_data3, by = "name"), 26)
```

```
##      name start_year      gym  weight_info pounds  treat
## 1   Nancy      2015      UFC starting_weight  180   cake
## 2   Nancy      2015      UFC starting_weight  180   cake
## 3   Nancy      2015      UFC starting_weight  180  choco
## 4   Nancy      2015      UFC starting_weight  180  choco
## 5   Nancy      2015      UFC starting_weight  180 icream
## 6   Nancy      2015      UFC starting_weight  180 icream
## 7  Sheldon      2014  24_HrFit starting_weight  220   cake
## 8  Sheldon      2014  24_HrFit starting_weight  220   cake
## 9  Sheldon      2014  24_HrFit starting_weight  220  choco
## 10 Sheldon      2014  24_HrFit starting_weight  220  choco
## 11 Sheldon      2014  24_HrFit starting_weight  220 icream
## 12 Sheldon      2014  24_HrFit starting_weight  220 icream
## 13 Samara      2016 Planet_Fitness starting_weight  170   cake
## 14 Samara      2016 Planet_Fitness starting_weight  170   cake
## 15 Samara      2016 Planet_Fitness starting_weight  170  choco
## 16 Samara      2016 Planet_Fitness starting_weight  170  choco
## 17 Samara      2016 Planet_Fitness starting_weight  170 icream
## 18 Samara      2016 Planet_Fitness starting_weight  170 icream
## 19 Brianna      2013  24_HrFit starting_weight  180   cake
## 20 Brianna      2013  24_HrFit starting_weight  180   cake
## 21 Brianna      2013  24_HrFit starting_weight  180  choco
## 22 Brianna      2013  24_HrFit starting_weight  180  choco
## 23 Brianna      2013  24_HrFit starting_weight  180 icream
## 24 Brianna      2013  24_HrFit starting_weight  180 icream
## 25   Nancy      2015      UFC      end_weight  155   cake
## 26   Nancy      2015      UFC      end_weight  155   cake
## serving miles
## 1      1  0.00
## 2      2  2.00
## 3      1  0.50
## 4      2  3.00
## 5      1  1.00
## 6      2  4.00
## 7      1  0.50
## 8      2  0.75
## 9      1  1.00
## 10     2  2.00
## 11     1  1.50
## 12     2  2.00
## 13     1  3.00
## 14     2  6.00
## 15     1  1.00
## 16     2  4.00
## 17     1  5.00
## 18     2  9.00
## 19     1  4.00
## 20     2  6.00
## 21     1  2.00
## 22     2  5.00
## 23     1  6.00
## 24     2 10.00
## 25     1  0.00
## 26     2  2.00
```

Let's consider another scenario. Suppose we want to join both data frames but that instead of being called "name" the names column of weight\_loss\_long was called "NAMES"?

```
# Change the name of the first column in weight_loss_long from "name" to "NAMES"

colnames(weight_loss_long)[1] <- "NAMES"
```

In order to be able to join the two data frames despite the fact that they have no identical column names as they did above, we use the following code:

```
head(inner_join(weight_loss_long, treats_data3, by = c("NAMES" = "name")), 50)
```

```
##      NAMES start_year      gym  weight_info pounds  treat
## 1   Nancy      2015      UFC starting_weight  180   cake
## 2   Nancy      2015      UFC starting_weight  180   cake
## 3   Nancy      2015      UFC starting_weight  180  choco
## 4   Nancy      2015      UFC starting_weight  180  choco
## 5   Nancy      2015      UFC starting_weight  180 icream
## 6   Nancy      2015      UFC starting_weight  180 icream
## 7  Sheldon      2014  24_HrFit starting_weight  220   cake
## 8  Sheldon      2014  24_HrFit starting_weight  220   cake
## 9  Sheldon      2014  24_HrFit starting_weight  220  choco
## 10 Sheldon      2014  24_HrFit starting_weight  220  choco
## 11 Sheldon      2014  24_HrFit starting_weight  220 icream
## 12 Sheldon      2014  24_HrFit starting_weight  220 icream
## 13 Samara      2016 Planet_Fitness starting_weight  170   cake
## 14 Samara      2016 Planet_Fitness starting_weight  170   cake
## 15 Samara      2016 Planet_Fitness starting_weight  170  choco
```

## 16	Samara	2016	Planet_Fitness	starting_weight	170	choco
## 17	Samara	2016	Planet_Fitness	starting_weight	170	icream
## 18	Samara	2016	Planet_Fitness	starting_weight	170	icream
## 19	Brianna	2013	24_HrFit	starting_weight	180	cake
## 20	Brianna	2013	24_HrFit	starting_weight	180	cake
## 21	Brianna	2013	24_HrFit	starting_weight	180	choco
## 22	Brianna	2013	24_HrFit	starting_weight	180	choco
## 23	Brianna	2013	24_HrFit	starting_weight	180	icream
## 24	Brianna	2013	24_HrFit	starting_weight	180	icream
## 25	Nancy	2015	UFC	end_weight	155	cake
## 26	Nancy	2015	UFC	end_weight	155	cake
## 27	Nancy	2015	UFC	end_weight	155	choco
## 28	Nancy	2015	UFC	end_weight	155	choco
## 29	Nancy	2015	UFC	end_weight	155	icream
## 30	Nancy	2015	UFC	end_weight	155	icream
## 31	Sheldon	2014	24_HrFit	end_weight	185	cake
## 32	Sheldon	2014	24_HrFit	end_weight	185	cake
## 33	Sheldon	2014	24_HrFit	end_weight	185	choco
## 34	Sheldon	2014	24_HrFit	end_weight	185	choco
## 35	Sheldon	2014	24_HrFit	end_weight	185	icream
## 36	Sheldon	2014	24_HrFit	end_weight	185	icream
## 37	Samara	2016	Planet_Fitness	end_weight	135	cake
## 38	Samara	2016	Planet_Fitness	end_weight	135	cake
## 39	Samara	2016	Planet_Fitness	end_weight	135	choco
## 40	Samara	2016	Planet_Fitness	end_weight	135	choco
## 41	Samara	2016	Planet_Fitness	end_weight	135	icream
## 42	Samara	2016	Planet_Fitness	end_weight	135	icream
## 43	Brianna	2013	24_HrFit	end_weight	150	cake
## 44	Brianna	2013	24_HrFit	end_weight	150	cake
## 45	Brianna	2013	24_HrFit	end_weight	150	choco
## 46	Brianna	2013	24_HrFit	end_weight	150	choco
## 47	Brianna	2013	24_HrFit	end_weight	150	icream
## 48	Brianna	2013	24_HrFit	end_weight	150	icream
## 49	Nancy	2015	UFC	goal_weight	150	cake
## 50	Nancy	2015	UFC	goal_weight	150	cake
##	serving	miles				
## 1	1	0.00				
## 2	2	2.00				
## 3	1	0.50				
## 4	2	3.00				
## 5	1	1.00				
## 6	2	4.00				
## 7	1	0.50				
## 8	2	0.75				
## 9	1	1.00				
## 10	2	2.00				
## 11	1	1.50				
## 12	2	2.00				
## 13	1	3.00				
## 14	2	6.00				
## 15	1	1.00				
## 16	2	4.00				
## 17	1	5.00				
## 18	2	9.00				
## 19	1	4.00				
## 20	2	6.00				
## 21	1	2.00				
## 22	2	5.00				
## 23	1	6.00				
## 24	2	10.00				
## 25	1	0.00				
## 26	2	2.00				
## 27	1	0.50				
## 28	2	3.00				
## 29	1	1.00				
## 30	2	4.00				
## 31	1	0.50				
## 32	2	0.75				
## 33	1	1.00				
## 34	2	2.00				
## 35	1	1.50				
## 36	2	2.00				
## 37	1	3.00				
## 38	2	6.00				
## 39	1	1.00				
## 40	2	4.00				
## 41	1	5.00				
## 42	2	9.00				
## 43	1	4.00				
## 44	2	6.00				
## 45	1	2.00				
## 46	2	5.00				
## 47	1	6.00				
## 48	2	10.00				
## 49	1	0.00				
## 50	2	2.00				

```
## 30      2      2.00
```

If you try to join the two data frames without including the code we set equal to the “by” parameter, you’ll get an error!

I’ll still be providing one more example using `inner_join()`, but this is a good place to introduce the `semi_join()` function.

```
semi_join(weight_loss_long, treats_data3, by = c("NAMES" = "name"))
```

```
##      NAMES start_year      gym  weight_info pounds
## 1   Nancy      2015      UFC starting_weight  180
## 2 Sheldon      2014    24_HrFit starting_weight  220
## 3   Samara      2016 Planet_Fitness starting_weight  170
## 4 Brianna      2013    24_HrFit starting_weight  180
## 5   Nancy      2015      UFC      end_weight  155
## 6 Sheldon      2014    24_HrFit      end_weight  185
## 7   Samara      2016 Planet_Fitness      end_weight  135
## 8 Brianna      2013    24_HrFit      end_weight  150
## 9   Nancy      2015      UFC    goal_weight  150
## 10 Sheldon      2014    24_HrFit    goal_weight  180
## 11 Samara      2016 Planet_Fitness    goal_weight  135
## 12 Brianna      2013    24_HrFit    goal_weight  140
```

The only difference between `semi_join()` and `inner_join()` is that `semi_join` only includes the columns of the data frame passed in as the first argument.

Just to solidify our understanding of `inner_join()`, as mentioned above I’ll provide one more example using this function.

Let’s suppose that the “NAMES” column of `weight_loss_long` only contains the names “Samara” and “Brianna.”

```
weight_loss_long <- weight_loss_long[-c(1,2,5,6,9,10), ]
weight_loss_long
```

```
##      NAMES start_year      gym  weight_info pounds
## 3   Samara      2016 Planet_Fitness starting_weight  170
## 4 Brianna      2013    24_HrFit starting_weight  180
## 7   Samara      2016 Planet_Fitness      end_weight  135
## 8 Brianna      2013    24_HrFit      end_weight  150
## 11 Samara      2016 Planet_Fitness    goal_weight  135
## 12 Brianna      2013    24_HrFit    goal_weight  140
```

Now we’re going to join `weight_loss_long` and `treats_data3` once more!

```
inner_join(treats_data3, weight_loss_long, by = c("name" = "NAMES"))
```

##	name	treat	serving	miles	start_year	gym	weight_info
## 1	Samara	cake	1	3	2016	Planet_Fitness	starting_weight
## 2	Samara	cake	1	3	2016	Planet_Fitness	end_weight
## 3	Samara	cake	1	3	2016	Planet_Fitness	goal_weight
## 4	Brianna	cake	1	4	2013	24_HrFit	starting_weight
## 5	Brianna	cake	1	4	2013	24_HrFit	end_weight
## 6	Brianna	cake	1	4	2013	24_HrFit	goal_weight
## 7	Samara	cake	2	6	2016	Planet_Fitness	starting_weight
## 8	Samara	cake	2	6	2016	Planet_Fitness	end_weight
## 9	Samara	cake	2	6	2016	Planet_Fitness	goal_weight
## 10	Brianna	cake	2	6	2013	24_HrFit	starting_weight
## 11	Brianna	cake	2	6	2013	24_HrFit	end_weight
## 12	Brianna	cake	2	6	2013	24_HrFit	goal_weight
## 13	Samara	choco	1	1	2016	Planet_Fitness	starting_weight
## 14	Samara	choco	1	1	2016	Planet_Fitness	end_weight
## 15	Samara	choco	1	1	2016	Planet_Fitness	goal_weight
## 16	Brianna	choco	1	2	2013	24_HrFit	starting_weight
## 17	Brianna	choco	1	2	2013	24_HrFit	end_weight
## 18	Brianna	choco	1	2	2013	24_HrFit	goal_weight
## 19	Samara	choco	2	4	2016	Planet_Fitness	starting_weight
## 20	Samara	choco	2	4	2016	Planet_Fitness	end_weight
## 21	Samara	choco	2	4	2016	Planet_Fitness	goal_weight
## 22	Brianna	choco	2	5	2013	24_HrFit	starting_weight
## 23	Brianna	choco	2	5	2013	24_HrFit	end_weight
## 24	Brianna	choco	2	5	2013	24_HrFit	goal_weight
## 25	Samara	icream	1	5	2016	Planet_Fitness	starting_weight
## 26	Samara	icream	1	5	2016	Planet_Fitness	end_weight
## 27	Samara	icream	1	5	2016	Planet_Fitness	goal_weight
## 28	Brianna	icream	1	6	2013	24_HrFit	starting_weight
## 29	Brianna	icream	1	6	2013	24_HrFit	end_weight
## 30	Brianna	icream	1	6	2013	24_HrFit	goal_weight
## 31	Samara	icream	2	9	2016	Planet_Fitness	starting_weight
## 32	Samara	icream	2	9	2016	Planet_Fitness	end_weight
## 33	Samara	icream	2	9	2016	Planet_Fitness	goal_weight
## 34	Brianna	icream	2	10	2013	24_HrFit	starting_weight
## 35	Brianna	icream	2	10	2013	24_HrFit	end_weight
## 36	Brianna	icream	2	10	2013	24_HrFit	goal_weight
##	pounds						
## 1	170						
## 2	135						
## 3	135						
## 4	180						
## 5	150						
## 6	140						
## 7	170						
## 8	135						
## 9	135						
## 10	180						
## 11	150						
## 12	140						
## 13	170						
## 14	135						
## 15	135						
## 16	180						
## 17	150						
## 18	140						
## 19	170						
## 20	135						
## 21	135						
## 22	180						
## 23	150						
## 24	140						
## 25	170						
## 26	135						
## 27	135						
## 28	180						
## 29	150						
## 30	140						
## 31	170						
## 32	135						
## 33	135						
## 34	180						
## 35	150						
## 36	140						

Great! We see that in the joined data frame we only have rows for “Samara” and “Brianna.”

Let’s also re-use semi\_join() after having removed the names “Sheldon” and “Nancy” from weight\_loss\_long

```
semi_j_df <- semi_join(treats_data3, weight_loss_long, by = c("name" = "NAMES"))

semi_j_df
```


```
##      name  treat serving miles
## 1  Samara   cake        1     3
## 2  Brianna  cake        1     4
## 3  Samara   cake        2     6
## 4  Brianna  cake        2     6
## 5  Samara  choco        1     1
## 6  Brianna  choco        1     2
## 7  Samara  choco        2     4
## 8  Brianna  choco        2     5
## 9  Samara icream        1     5
##10  Brianna icream        1     6
##11  Samara icream        2     9
##12  Brianna icream        2    10
```

Here we can see that `semi_join()` returns a data frame that is almost identical to `treats_data3`, except for the fact that it has excluded the rows corresponding to the names "Sheldon" and "Nancy" since those names don't exist in `weight_loss_long`.

Let's create another barchart comparing the total number of miles that Samara and Brianna will have run after they have, at some point, consumed every `treat_serving` combination.

```
# Compare total number of miles run by Brianna and Samara after they have eaten every treat-serving combo

ggplot(semi_j_df, aes(x = name, y = miles, fill = name)) + geom_bar(stat = 'identity') + labs(x = "Names", y = "Total Number of Miles")
```



Brianna clearly ends up running more!

Ok, let's see how using `left_join()` gives us a different output than that produced by `inner_join()`!

```
# Using left_join() with treats_data3 as the x parameter (first argument) and weight_loss_long as the y parameter (second argument)

left_join1 <- left_join(treats_data3, weight_loss_long, by = c("name" = "NAMES"))

head(left_join1, 20)
```

```
##      name treat serving miles start_year      gym      weight_info
## 1  Samara  cake      1   3.00      2016 Planet_Fitness starting_weight
## 2  Samara  cake      1   3.00      2016 Planet_Fitness      end_weight
## 3  Samara  cake      1   3.00      2016 Planet_Fitness    goal_weight
## 4   Nancy  cake      1   0.00        NA          <NA>          <NA>
## 5 Sheldon  cake      1   0.50        NA          <NA>          <NA>
## 6 Brianna  cake      1   4.00      2013      24_HrFit starting_weight
## 7 Brianna  cake      1   4.00      2013      24_HrFit      end_weight
## 8 Brianna  cake      1   4.00      2013      24_HrFit    goal_weight
## 9  Samara  cake      2   6.00      2016 Planet_Fitness starting_weight
## 10 Samara  cake      2   6.00      2016 Planet_Fitness      end_weight
## 11 Samara  cake      2   6.00      2016 Planet_Fitness    goal_weight
## 12 Nancy  cake      2   2.00        NA          <NA>          <NA>
## 13 Sheldon  cake      2   0.75        NA          <NA>          <NA>
## 14 Brianna  cake      2   6.00      2013      24_HrFit starting_weight
## 15 Brianna  cake      2   6.00      2013      24_HrFit      end_weight
## 16 Brianna  cake      2   6.00      2013      24_HrFit    goal_weight
## 17 Samara  choco      1   1.00      2016 Planet_Fitness starting_weight
## 18 Samara  choco      1   1.00      2016 Planet_Fitness      end_weight
## 19 Samara  choco      1   1.00      2016 Planet_Fitness    goal_weight
## 20 Nancy  choco      1   0.50        NA          <NA>          <NA>
##      pounds
## 1      170
## 2      135
## 3      135
## 4       NA
## 5       NA
## 6      180
## 7      150
## 8      140
## 9      170
## 10     135
## 11     135
## 12      NA
## 13      NA
## 14      180
## 15      150
## 16      140
## 17      170
## 18      135
## 19      135
## 20      NA
```

```
# Using left_join() with weight_loss_long as the x parameter (first argument) and treats_data3 as the y parameter (second argument)
```

```
left_join2 <- left_join(weight_loss_long, treats_data3, by = c("NAMES" = "name"))

head(left_join2, 24)
```

```
##      NAMES start_year      gym      weight_info pounds treat
## 1  Samara      2016 Planet_Fitness starting_weight    170  cake
## 2  Samara      2016 Planet_Fitness starting_weight    170  cake
## 3  Samara      2016 Planet_Fitness starting_weight    170 choco
## 4  Samara      2016 Planet_Fitness starting_weight    170 choco
## 5  Samara      2016 Planet_Fitness starting_weight    170 icream
## 6  Samara      2016 Planet_Fitness starting_weight    170 icream
## 7  Brianna     2013      24_HrFit starting_weight    180  cake
## 8  Brianna     2013      24_HrFit starting_weight    180  cake
## 9  Brianna     2013      24_HrFit starting_weight    180 choco
## 10 Brianna     2013      24_HrFit starting_weight    180 choco
## 11 Brianna     2013      24_HrFit starting_weight    180 icream
## 12 Brianna     2013      24_HrFit starting_weight    180 icream
## 13 Samara      2016 Planet_Fitness      end_weight    135  cake
## 14 Samara      2016 Planet_Fitness      end_weight    135  cake
## 15 Samara      2016 Planet_Fitness      end_weight    135 choco
## 16 Samara      2016 Planet_Fitness      end_weight    135 choco
## 17 Samara      2016 Planet_Fitness      end_weight    135 icream
## 18 Samara      2016 Planet_Fitness      end_weight    135 icream
## 19 Brianna     2013      24_HrFit      end_weight    150  cake
## 20 Brianna     2013      24_HrFit      end_weight    150  cake
## 21 Brianna     2013      24_HrFit      end_weight    150 choco
## 22 Brianna     2013      24_HrFit      end_weight    150 choco
## 23 Brianna     2013      24_HrFit      end_weight    150 icream
## 24 Brianna     2013      24_HrFit      end_weight    150 icream
##      serving miles
## 1         1      3
## 2         2      6
## 3         1      1
## 4         2      4
## 5         1      5
## 6         2      9
## 7         1      4
## 8         2      6
## 9         1      2
## 10        2      5
## 11        1      6
## 12        2     10
## 13        1      3
## 14        2      6
## 15        1      1
## 16        2      4
## 17        1      5
## 18        2      9
## 19        1      4
## 20        2      6
## 21        1      2
## 22        2      5
## 23        1      6
## 24        2     10
```

Ok, so there are a few things going on here. To begin with, we see that `left_join1` and `left_join2` do *not* look the same (contain the exact same contents)! This is because `left_join()` returns every row from the data frame passed in as the first argument as well as every column from *both* data frames. In `left_join1` we are returning all of the rows from `treats_data3` as well as all of the columns from `treats_data3` *and* `weight_loss_long`. This is why we see NA values in our data frame. There are rows in `treats_data3` that are not in `weight_loss_long`, namely those for “Nancy” and “Sheldon”. Therefore, we get NA in the columns from `weight_loss_long` for both of those names. We do not, on the other hand, get any NA values in `left_join2`, because for every name in `weight_loss_long` there is a corresponding name in `treats_data3` and thus input available for that name in every column of `treats_data3`.

Let’s move onto `right_join()`! The only difference between `right_join()` and `left_join()` is that `right_join()` will return all of the rows from the data frame that is passed in as the second argument of `right_join()`, as opposed to the first (hence the “right” join).

Therefore, if we want to get a data frame that is identical to `left_join1` using the `right_join()` function, all we have to do is the following:

```
right_join1 <- right_join(weight_loss_long, treats_data3, by = c("NAMES" = "name"))

head(right_join1, 20)
```



```
##      NAMES start_year      gym      weight_info pounds treat serving
## 1  Samara      2016 Planet_Fitness starting_weight    170 cake      1
## 2  Samara      2016 Planet_Fitness      end_weight    135 cake      1
## 3  Samara      2016 Planet_Fitness      goal_weight    135 cake      1
## 4   Nancy        NA          <NA>          <NA>      NA cake      1
## 5 Sheldon        NA          <NA>          <NA>      NA cake      1
## 6 Brianna      2013      24_HrFit starting_weight    180 cake      1
## 7 Brianna      2013      24_HrFit      end_weight    150 cake      1
## 8 Brianna      2013      24_HrFit      goal_weight    140 cake      1
## 9  Samara      2016 Planet_Fitness starting_weight    170 cake      2
## 10 Samara      2016 Planet_Fitness      end_weight    135 cake      2
## 11 Samara      2016 Planet_Fitness      goal_weight    135 cake      2
## 12 Nancy        NA          <NA>          <NA>      NA cake      2
## 13 Sheldon        NA          <NA>          <NA>      NA cake      2
## 14 Brianna      2013      24_HrFit starting_weight    180 cake      2
## 15 Brianna      2013      24_HrFit      end_weight    150 cake      2
## 16 Brianna      2013      24_HrFit      goal_weight    140 cake      2
## 17 Samara      2016 Planet_Fitness starting_weight    170 choco     1
## 18 Samara      2016 Planet_Fitness      end_weight    135 choco     1
## 19 Samara      2016 Planet_Fitness      goal_weight    135 choco     1
## 20 Nancy        NA          <NA>          <NA>      NA choco     1
##      miles
## 1      3.00
## 2      3.00
## 3      3.00
## 4      0.00
## 5      0.50
## 6      4.00
## 7      4.00
## 8      4.00
## 9      6.00
## 10     6.00
## 11     6.00
## 12     2.00
## 13     0.75
## 14     6.00
## 15     6.00
## 16     6.00
## 17     1.00
## 18     1.00
## 19     1.00
## 20     0.50
```

Now to briefly discuss `full_join()`! As the name implies, `full_join()` simply returns a data frame that contains all of the rows and columns from the data frames passed in as the first and second arguments of the function.

```
full_join1 <- full_join(treats_data3, weight_loss_long, by = c("name" = "NAMES"))

head(full_join1, 30)
```

```
##      name treat serving miles start_year      gym      weight_info
## 1  Samara  cake      1  3.00      2016 Planet_Fitness starting_weight
## 2  Samara  cake      1  3.00      2016 Planet_Fitness      end_weight
## 3  Samara  cake      1  3.00      2016 Planet_Fitness      goal_weight
## 4   Nancy  cake      1  0.00        NA          <NA>          <NA>
## 5  Sheldon  cake      1  0.50        NA          <NA>          <NA>
## 6  Brianna  cake      1  4.00      2013      24_HrFit starting_weight
## 7  Brianna  cake      1  4.00      2013      24_HrFit      end_weight
## 8  Brianna  cake      1  4.00      2013      24_HrFit      goal_weight
## 9  Samara  cake      2  6.00      2016 Planet_Fitness starting_weight
## 10 Samara  cake      2  6.00      2016 Planet_Fitness      end_weight
## 11 Samara  cake      2  6.00      2016 Planet_Fitness      goal_weight
## 12  Nancy  cake      2  2.00        NA          <NA>          <NA>
## 13 Sheldon  cake      2  0.75        NA          <NA>          <NA>
## 14 Brianna  cake      2  6.00      2013      24_HrFit starting_weight
## 15 Brianna  cake      2  6.00      2013      24_HrFit      end_weight
## 16 Brianna  cake      2  6.00      2013      24_HrFit      goal_weight
## 17 Samara  choco      1  1.00      2016 Planet_Fitness starting_weight
## 18 Samara  choco      1  1.00      2016 Planet_Fitness      end_weight
## 19 Samara  choco      1  1.00      2016 Planet_Fitness      goal_weight
## 20  Nancy  choco      1  0.50        NA          <NA>          <NA>
## 21 Sheldon  choco      1  1.00        NA          <NA>          <NA>
## 22 Brianna  choco      1  2.00      2013      24_HrFit starting_weight
## 23 Brianna  choco      1  2.00      2013      24_HrFit      end_weight
## 24 Brianna  choco      1  2.00      2013      24_HrFit      goal_weight
## 25 Samara  choco      2  4.00      2016 Planet_Fitness starting_weight
## 26 Samara  choco      2  4.00      2016 Planet_Fitness      end_weight
## 27 Samara  choco      2  4.00      2016 Planet_Fitness      goal_weight
## 28  Nancy  choco      2  3.00        NA          <NA>          <NA>
## 29 Sheldon  choco      2  2.00        NA          <NA>          <NA>
## 30 Brianna  choco      2  5.00      2013      24_HrFit starting_weight
##      pounds
## 1      170
## 2      135
## 3      135
## 4      NA
## 5      NA
## 6      180
## 7      150
## 8      140
## 9      170
## 10     135
## 11     135
## 12     NA
## 13     NA
## 14     180
## 15     150
## 16     140
## 17     170
## 18     135
## 19     135
## 20     NA
## 21     NA
## 22     180
## 23     150
## 24     140
## 25     170
## 26     135
## 27     135
## 28     NA
## 29     NA
## 30     180
```

In our case, both `full_join1` and `left_join1` are the same!

### Removing NA Values

Let's suppose that we want to eliminate the rows containing the NA values in `full_join1`.

In our case, we know that doing so will output a data frame with the same contents as `left_join2` (where we have data for "Samara" and "Brianna"—but *do not* have data for "Nancy" and "Sheldon"—for all of the columns in both `treats_data3` and `weight_loss_long`). For the sake of demonstrating the use of `complete.cases()` and `na.omit()`, however, I'll go ahead and use `full_join1`.

I'll start off by saying that `na.omit()` is significantly easier to use, as you only have to pass a data frame into the function in order to remove the rows containing NA values. Because `na.omit()` is pretty straightforward, I'll start off by showing how to use `complete.cases()`

When using `complete.cases()` we will be selecting the columns that contain NA values.

```
# Removing rows containing NA values from full_join1 using complete.cases()

head(full_join1[complete.cases(full_join1[,5:8]), ], 25)
```

```
##      name treat serving miles start_year      gym      weight_info
## 1  Samara  cake      1      3      2016 Planet_Fitness starting_weight
## 2  Samara  cake      1      3      2016 Planet_Fitness      end_weight
## 3  Samara  cake      1      3      2016 Planet_Fitness      goal_weight
## 6  Brianna cake      1      4      2013      24_HrFit starting_weight
## 7  Brianna cake      1      4      2013      24_HrFit      end_weight
## 8  Brianna cake      1      4      2013      24_HrFit      goal_weight
## 9  Samara  cake      2      6      2016 Planet_Fitness starting_weight
## 10 Samara  cake      2      6      2016 Planet_Fitness      end_weight
## 11 Samara  cake      2      6      2016 Planet_Fitness      goal_weight
## 14 Brianna cake      2      6      2013      24_HrFit starting_weight
## 15 Brianna cake      2      6      2013      24_HrFit      end_weight
## 16 Brianna cake      2      6      2013      24_HrFit      goal_weight
## 17 Samara choco      1      1      2016 Planet_Fitness starting_weight
## 18 Samara choco      1      1      2016 Planet_Fitness      end_weight
## 19 Samara choco      1      1      2016 Planet_Fitness      goal_weight
## 22 Brianna choco      1      2      2013      24_HrFit starting_weight
## 23 Brianna choco      1      2      2013      24_HrFit      end_weight
## 24 Brianna choco      1      2      2013      24_HrFit      goal_weight
## 25 Samara  choco      2      4      2016 Planet_Fitness starting_weight
## 26 Samara  choco      2      4      2016 Planet_Fitness      end_weight
## 27 Samara  choco      2      4      2016 Planet_Fitness      goal_weight
## 30 Brianna choco      2      5      2013      24_HrFit starting_weight
## 31 Brianna choco      2      5      2013      24_HrFit      end_weight
## 32 Brianna choco      2      5      2013      24_HrFit      goal_weight
## 33 Samara icream      1      5      2016 Planet_Fitness starting_weight
##      pounds
## 1      170
## 2      135
## 3      135
## 6      180
## 7      150
## 8      140
## 9      170
## 10     135
## 11     135
## 14     180
## 15     150
## 16     140
## 17     170
## 18     135
## 19     135
## 22     180
## 23     150
## 24     140
## 25     170
## 26     135
## 27     135
## 30     180
## 31     150
## 32     140
## 33     170
```

Awesome! No more rows with NA values.

Ok, so what exactly is this code doing? To begin with, we see that we subset `full_join1` in order to get all of the rows for columns 5 through 8. We only want columns 5:8, because these are the columns that contain NA values. So we see that the data frame being passed into `complete.cases()` is the following:

```
full_join1[,5:8]
```

##	start_year	gym	weight_info	pounds
## 1	2016	Planet_Fitness	starting_weight	170
## 2	2016	Planet_Fitness	end_weight	135
## 3	2016	Planet_Fitness	goal_weight	135
## 4	NA	<NA>	<NA>	NA
## 5	NA	<NA>	<NA>	NA
## 6	2013	24_HrFit	starting_weight	180
## 7	2013	24_HrFit	end_weight	150
## 8	2013	24_HrFit	goal_weight	140
## 9	2016	Planet_Fitness	starting_weight	170
## 10	2016	Planet_Fitness	end_weight	135
## 11	2016	Planet_Fitness	goal_weight	135
## 12	NA	<NA>	<NA>	NA
## 13	NA	<NA>	<NA>	NA
## 14	2013	24_HrFit	starting_weight	180
## 15	2013	24_HrFit	end_weight	150
## 16	2013	24_HrFit	goal_weight	140
## 17	2016	Planet_Fitness	starting_weight	170
## 18	2016	Planet_Fitness	end_weight	135
## 19	2016	Planet_Fitness	goal_weight	135
## 20	NA	<NA>	<NA>	NA
## 21	NA	<NA>	<NA>	NA
## 22	2013	24_HrFit	starting_weight	180
## 23	2013	24_HrFit	end_weight	150
## 24	2013	24_HrFit	goal_weight	140
## 25	2016	Planet_Fitness	starting_weight	170
## 26	2016	Planet_Fitness	end_weight	135
## 27	2016	Planet_Fitness	goal_weight	135
## 28	NA	<NA>	<NA>	NA
## 29	NA	<NA>	<NA>	NA
## 30	2013	24_HrFit	starting_weight	180
## 31	2013	24_HrFit	end_weight	150
## 32	2013	24_HrFit	goal_weight	140
## 33	2016	Planet_Fitness	starting_weight	170
## 34	2016	Planet_Fitness	end_weight	135
## 35	2016	Planet_Fitness	goal_weight	135
## 36	NA	<NA>	<NA>	NA
## 37	NA	<NA>	<NA>	NA
## 38	2013	24_HrFit	starting_weight	180
## 39	2013	24_HrFit	end_weight	150
## 40	2013	24_HrFit	goal_weight	140
## 41	2016	Planet_Fitness	starting_weight	170
## 42	2016	Planet_Fitness	end_weight	135
## 43	2016	Planet_Fitness	goal_weight	135
## 44	NA	<NA>	<NA>	NA
## 45	NA	<NA>	<NA>	NA
## 46	2013	24_HrFit	starting_weight	180
## 47	2013	24_HrFit	end_weight	150
## 48	2013	24_HrFit	goal_weight	140

What `complete.cases()` does is take this data frame and output a logical vector specifying which rows have NA values for the given columns. FALSE indicates that the row has NA values. Let's see what this logical vector looks like.

```
# Logical vector output by complete.cases()

complete_cases <- complete.cases(full_join1[,5:8])

complete_cases
```

```
## [1] TRUE TRUE TRUE FALSE FALSE TRUE TRUE TRUE TRUE TRUE TRUE
## [12] FALSE FALSE TRUE TRUE TRUE TRUE FALSE FALSE TRUE TRUE TRUE TRUE
## [23] TRUE TRUE TRUE TRUE TRUE TRUE FALSE FALSE TRUE TRUE TRUE TRUE
## [34] TRUE TRUE FALSE FALSE TRUE TRUE TRUE TRUE TRUE TRUE FALSE
## [45] FALSE TRUE TRUE TRUE
```

Here we can see that the indices of the FALSE values in the vector `complete_cases` correspond to the rows that contain NA values in `full_join1`. For example, we see that rows 4 and 5 of `full_join1` contain NA values and that the 4th and 5th indices of `complete_cases` are FALSE values. This is why when we use `complete_cases` to subset `full_join1` we get only the rows without any NA values.

One can also use `na.omit()` to get the same output.

```
head(na.omit(full_join1), 25)
```

```
##      name treat serving miles start_year      gym      weight_info
## 1   Samara  cake        1      3        2016 Planet_Fitness starting_weight
## 2   Samara  cake        1      3        2016 Planet_Fitness      end_weight
## 3   Samara  cake        1      3        2016 Planet_Fitness    goal_weight
## 6   Brianna cake        1      4        2013      24_HrFit starting_weight
## 7   Brianna cake        1      4        2013      24_HrFit      end_weight
## 8   Brianna cake        1      4        2013      24_HrFit    goal_weight
## 9   Samara  cake        2      6        2016 Planet_Fitness starting_weight
## 10  Samara  cake        2      6        2016 Planet_Fitness      end_weight
## 11  Samara  cake        2      6        2016 Planet_Fitness    goal_weight
## 14  Brianna cake        2      6        2013      24_HrFit starting_weight
## 15  Brianna cake        2      6        2013      24_HrFit      end_weight
## 16  Brianna cake        2      6        2013      24_HrFit    goal_weight
## 17  Samara  choco       1      1        2016 Planet_Fitness starting_weight
## 18  Samara  choco       1      1        2016 Planet_Fitness      end_weight
## 19  Samara  choco       1      1        2016 Planet_Fitness    goal_weight
## 22  Brianna choco       1      2        2013      24_HrFit starting_weight
## 23  Brianna choco       1      2        2013      24_HrFit      end_weight
## 24  Brianna choco       1      2        2013      24_HrFit    goal_weight
## 25  Samara  choco       2      4        2016 Planet_Fitness starting_weight
## 26  Samara  choco       2      4        2016 Planet_Fitness      end_weight
## 27  Samara  choco       2      4        2016 Planet_Fitness    goal_weight
## 30  Brianna choco       2      5        2013      24_HrFit starting_weight
## 31  Brianna choco       2      5        2013      24_HrFit      end_weight
## 32  Brianna choco       2      5        2013      24_HrFit    goal_weight
## 33  Samara  icream      1      5        2016 Planet_Fitness starting_weight
##      pounds
## 1      170
## 2      135
## 3      135
## 6      180
## 7      150
## 8      140
## 9      170
## 10     135
## 11     135
## 14     180
## 15     150
## 16     140
## 17     170
## 18     135
## 19     135
## 22     180
## 23     150
## 24     140
## 25     170
## 26     135
## 27     135
## 30     180
## 31     150
## 32     140
## 33     170
```

This is obviously much simpler, but it's still worth knowing how to use `complete.cases()`!

### Exploring a Few `apply()` Functions

I'll now be moving on to talk about the family of `apply()` functions. The `apply()` function itself has popped up a few times in Stat 133, however, I'd like to discuss this function a bit more and introduce the following two `apply()` functions:

- `mapply()`
- `lapply()`

I'd like to start off with a simple `apply()` example in order to demonstrate how one can use `apply()` to perform a function either by row or by column. It is the "MARGIN" parameter of the `apply()` function that allows for this specification.

```
# Let's generate a random data frame

apply_df <- data.frame(replicate(5, sample(1:10, 6, rep = TRUE)))

apply_df
```

```
##      X1 X2 X3 X4 X5
## 1    3  4  1  2  1
## 2    1  8  8  2  9
## 3    9  1  2  3  7
## 4    3  3  8  1 10
## 5    8  2  4  7  6
## 6    1  2  2  7  4
```

```
# Multiplying by row using apply()

apply(apply_df, 1, prod)
```

```
## [1] 24 1152 378 720 2688 112
```

```
# Multiplying by column using apply()

apply(apply_df, 2, prod)
```

```
##      X1      X2      X3      X4      X5
## 648    384   1024    588 15120
```

Great! So we see that if we want to apply a function by *row* we set “MARGIN” equal to **1**, and if we want to apply a function by *column* we set “MARGIN” equal to **2**.

Ok, let's see how `mapply()` works. Keep in mind that `mapply()` allows a function to be applied to several vectors or lists.

```
# Here is one example using mapply()
mapply_df <- as.data.frame(mapply(rep, 2:5, 3))

mapply_df
```

```
##      V1 V2 V3 V4
## 1  2  3  4  5
## 2  2  3  4  5
## 3  2  3  4  5
```

Had we not used `mapply()`, we would have had to do something like this:

```
# Generating mapply_df without using mapply()

as.data.frame(matrix(c(rep(2, 3), rep(3,3), rep(4,3), rep(5,3)), 3, 4))
```

```
##      V1 V2 V3 V4
## 1  2  3  4  5
## 2  2  3  4  5
## 3  2  3  4  5
```

Clearly less efficient! Ok, let's go over one more example using `mapply()`.

```
list_1 <- list(1:4, 6:9, 10:13)
list_2 <- list(4:7, 12:15, 10:16)

mapply(prod, list_1, list_2)
```

```
## [1] 20160 99066240 989404416000
```

So here we see that the corresponding elements of each list are being multiplied by each other. In other words 1:4 from `list_1` is being multiplied by 4:7 of `list_2`. The numbers within each of those vectors are also multiplied by one another before being multiplied by the vector of the other list, so when we multiply 1:4 by 4:7, we are actually multiplying the product of 1, 2, 3, and 4 by the product of 4, 5, 6, and 7.

What if we wanted to use `mapply()` on only part of each list as opposed to the whole list? We could do the following:

```
# Use mapply() to multiply the second element of both lists by one another

mapply(prod, list_1[2], list_2[2])
```

```
## [1] 99066240
```

It is important to note that the lists must be of the same length if we wish to apply a function to them in the way that we do with `mapply()`. In other words, if `list_2` contained only two elements as opposed to three, trying to run the code `mapply(prod, list_1, list_2)` would **not** work.

Ok, now let's go over how to use `lapply()`! This function is very similar to `apply()`, the main difference being that `lapply()` outputs a list. I'll show one very simply example using `lapply()` and then one that is a bit more involved.

```
# Very simple application of lapply()
lapply_df <- data.frame(replicate(5, sample(1:10, 4, rep = TRUE)))

lapply_df
```

```
##      X1 X2 X3 X4 X5
## 1  8  4  3  4  5
## 2  5  3  4  8  2
## 3  5  8  6  7  1
## 4  6  2  8  3  8
```

```
lapply(lapply_df, median)
```

```
## $X1
## [1] 5.5
##
## $X2
## [1] 3.5
##
## $X3
## [1] 5
##
## $X4
## [1] 5.5
##
## $X5
## [1] 3.5
```

```
# Another example using lapply()

# First generate three random data frames

df_1 <- data.frame(replicate(5, sample(4:20, 4, rep = TRUE)))

df_2 <- data.frame(replicate(5, sample(8:30, 4, rep = TRUE)))

df_3 <- data.frame(replicate(5, sample(40:50, 4, rep = TRUE)))

# Combine these data frames into a list
df_list <- list(df_1, df_2, df_3)
```

Now we'll actually use `lapply()`!

```
# Selecting rows 3 and 4 from each data frame of df_list

lapply(df_list, "[", 3:4, )
```

```
## [[1]]
##   X1 X2 X3 X4 X5
## 3 12 20  4 15 13
## 4  5 12 12  4 16
##
## [[2]]
##   X1 X2 X3 X4 X5
## 3 14 25 30 12 25
## 4 26 13 12 16 16
##
## [[3]]
##   X1 X2 X3 X4 X5
## 3 41 44 45 50 41
## 4 40 48 46 50 46
```

```
# Selecting columns 3 and 4 from each data frame of df_list

lapply(df_list, "[", ,3:4)
```

```
## [[1]]
##   X3 X4
## 1 20 15
## 2 15 19
## 3  4 15
## 4 12  4
##
## [[2]]
##   X3 X4
## 1 19 17
## 2 30  9
## 3 30 12
## 4 12 16
##
## [[3]]
##   X3 X4
## 1 47 40
## 2 49 44
## 3 45 50
## 4 46 50
```

The output is still a list!

## Take-Home Message

This post has gone over several types of data frame manipulation. We began by highlighting the relationship between `aggregate()` and the “dplyr” functions `summarise()` and `group_by()`. The `aggregate()` function was also explored in greater depth than was done in Stat 133. Afterwards, we went over how to make an “untidy” data frame “tidy” as well how to switch between “wide” and “long” formats. Lastly, we discussed the

various **join()** functions provided by “dplyr” as well as a few of the functions within the **apply()** family. Overall, one can see that data frame manipulation can involve altering the **contents** of a data frame as well as its *structure*.



You got through the entire post!

## ***References***

I used the sites listed below when creating this post. I would like to give credit to David Kun, Sharon Machlis, Philippe Marchand, Karlijn Willems, Carlo Fanara, and Neil Sanders for inspiring the examples in this post.

[Site 1](#)

[Site 2](#)

[Site 3](#)

[Site 4](#)

[Site 5](#)

[Site 6](#)

[Site 7](#)

[Site 8](#)

[Site 9](#)