

Post01: Data Visualization with dplyr

Tianshu Zhao

2017.10.30

Introduction

"dplyr" is a package used for data manipulation in RStudio, written and maintained by *Hadley Wickham* (my 133 God). It provides some powerful, easy-to-use functions that are extremely useful for data exploration and analysis. As a Statistics and applied math senior, I found "dplyr" is my biggest time-saver. From what I can tell, you may not need to know the classification of the various machine learning and clustering algorithm, but if you use R to analysis data and don't know how to use "dplyr", then it is a huge pity. In this post, I want to focus on the things we have learned from class and it's further extension. Although we have seen most of the function in class, but we are not deeply familiar with it so far.



Ship code to data,
Functional Programming



Image

Background

"dplyr" is the upgrade version of "plyr" package, the data can be easily filtered, deformed, summarized, grouped and piped using "dplyr" to perform data processing. It covers 90% of user demand in RStudio.

In data analysis process, the original data set is often uncleaned, unsorted and non-transformed. The Common work for mining and transforming data mainly includes: specific analysis on the record of the selection of variables, meet the conditions of filter, sort by one or several variables, process of original variable and generate a new variable, summarize the data and grouping elements, such as calculating the average and standard deviation of each group.

According to the website of RStudio, the writer of **dplyr**—*Hadley Wickham* (also the writer of **ggplot2** package), he claims himself as "*a grammar of data manipulation*". He further isolated the **ddply()** functions in the "plyr" package and focused on accepting the **data frame** object, it has greatly increased the speed of data manipulation, and providing a more robust interface with other database objects.

This project tries to briefly introduce some basic and common functions of **dplyr()** package. It mainly includes:

- Variable filter function—**select**
- Character selection function—**filter**
- Order arrangement function—**arrange**
- Deformation(calculation) function —**mutate**
- Summary function—**summarize**
- grouping function—**group_by**
- Multi-step operation concatenations—**%>%**
- Simple random sample function—**sample_n, sample_frac**

Cognitive process:

1. Take the **ydat** dataset, *then*
2. **filter()** for genes in the leucine biosynthesis pathway, *then*
3. **group_by()** the limiting nutrient, *then*
4. **summarize()** to correlate rate and expression, *then*
5. **mutate()** to round *r* to two digits, *then*
6. **arrange()** by rounded correlation coefficients

The old way:

```
arrange(
  mutate(
    summarize(
      group_by(
        filter(ydat, bp=="leucine biosynthesis"),
        nutrient),
      r=cor(rate, expression)),
    r=round(r, 2)),
  r)
```

The dplyr way:

```
ydat %>%
  filter(bp=="leucine biosynthesis") %>%
  group_by(nutrient) %>%
  summarize(r=cor(rate, expression)) %>%
  mutate(r=round(r,2)) %>%
  arrange(r)
```

Image

First, we install the **dplyr** package and use the default **nycflight13** data as an example.

```
library(dplyr)
```

```
## Warning: package 'dplyr' was built under R version 3.4.2
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(nycflights13)
```

```
## Warning: package 'nycflights13' was built under R version 3.4.2
```

Please ignore the warning, my R is the newest version, I really have no idea how to fix it

Data overview

Data can sometimes have a lot of rows, and if you print it all at once, it will take a lot of time, you can not see the name of each row as well. So the R language provides us with a **head()** function, also in **dplyr**, there is a implements similar **tbl_df()** function, displaying the following results. You can see the years, months, days, departure time, schedule arrival time, actual arrive time, delay time, etc.

Output shows that, "flight" is a data frame, and it contains 336776 rows of data, 19 variables. **head()** function only display the first 6 rows.

```
head(flights)
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517             515           2     830
## 2  2013     1     1     533             529           4     850
## 3  2013     1     1     542             540           2     923
## 4  2013     1     1     544             545          -1    1004
## 5  2013     1     1     554             600          -6     812
## 6  2013     1     1     554             558          -4     740
## # ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

```
# source : local data frame [6 x 19]
tbl_df(flights)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517             515           2     830
## 2  2013     1     1     533             529           4     850
## 3  2013     1     1     542             540           2     923
## 4  2013     1     1     544             545          -1    1004
## 5  2013     1     1     554             600          -6     812
## 6  2013     1     1     554             558          -4     740
## 7  2013     1     1     555             600          -5     913
## 8  2013     1     1     557             600          -3     709
## 9  2013     1     1     557             600          -3     838
## 10 2013     1     1     558             600          -2     753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# source : local data frame [336, 776 x 19]
dim(flights)
```

```
## [1] 336776      19
```

```
class(flights)
```

```
## [1] "tbl_df"      "tbl"        "data.frame"
```

Variable filter function

The feature was previously also implemented using indexes, and dplyr uses the select function to make filtering more convenient.

```
select(flights, year, month, day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

```
select(flights, year:day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

The 2 codes above represents select the first 3 columns of data (year,month,day) Besides, we could also use the **distinct** function to filter duplicate rows according to the values of a column.

```
distinct(select(flights, origin,dest))
```

```
## # A tibble: 224 x 2
##   origin dest
##   <chr> <chr>
## 1   EWR   IAH
## 2   LGA   IAH
## 3   JFK   MIA
## 4   JFK   BQN
## 5   LGA   ATL
## 6   EWR   ORD
## 7   EWR   FLL
## 8   LGA   IAD
## 9   JFK   MCO
## 10  LGA   ORD
## # ... with 214 more rows
```

The code above represents the data of all rows that are not identical to the destination combinations.

Character selection function

filter() function provides a basic data screening. In the past, we used data.frame to screen the data in the index. For example, we wanted to find the data of January 1 to use Rcode:

```
flights[flights$month == 1 & flights$day == 1, ]
```

```
## # A tibble: 842 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     1     517             515         2     830
## 2  2013     1     1     533             529         4     850
## 3  2013     1     1     542             540         2     923
## 4  2013     1     1     544             545        -1    1004
## 5  2013     1     1     554             600        -6     812
## 6  2013     1     1     554             558        -4     740
## 7  2013     1     1     555             600        -5     913
## 8  2013     1     1     557             600        -3     709
## 9  2013     1     1     557             600        -3     838
## 10 2013     1     1     558             600        -2     753
## # ... with 832 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

In **dplyr** it provides a filter function that makes it easier to implement the above functionality:

```
filter(flights, month == 1, day ==1)
```

```
## # A tibble: 842 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517             515           2     830
## 2  2013     1     1     533             529           4     850
## 3  2013     1     1     542             540           2     923
## 4  2013     1     1     544             545          -1    1004
## 5  2013     1     1     554             600          -6     812
## 6  2013     1     1     554             558          -4     740
## 7  2013     1     1     555             600          -5     913
## 8  2013     1     1     557             600          -3     709
## 9  2013     1     1     557             600          -3     838
## 10 2013     1     1     558             600          -2     753
## # ... with 832 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

For example, the statement will often be use :

```
filter(flights, month == 1 | month == 2)
```

```
## # A tibble: 51,955 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517             515           2     830
## 2  2013     1     1     533             529           4     850
## 3  2013     1     1     542             540           2     923
## 4  2013     1     1     544             545          -1    1004
## 5  2013     1     1     554             600          -6     812
## 6  2013     1     1     554             558          -4     740
## 7  2013     1     1     555             600          -5     913
## 8  2013     1     1     557             600          -3     709
## 9  2013     1     1     557             600          -3     838
## 10 2013     1     1     558             600          -2     753
## # ... with 51,945 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

Variable filter function arrange

Besides `filter()` function, `dplyr` also provides a `arrange()` function that could help user to reorder the rows.

```
arrange(flights, year, month, day)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517             515           2     830
## 2  2013     1     1     533             529           4     850
## 3  2013     1     1     542             540           2     923
## 4  2013     1     1     544             545          -1    1004
## 5  2013     1     1     554             600          -6     812
## 6  2013     1     1     554             558          -4     740
## 7  2013     1     1     555             600          -5     913
## 8  2013     1     1     557             600          -3     709
## 9  2013     1     1     557             600          -3     838
## 10 2013     1     1     558             600          -2     753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

we can also use the `desc` keyword to reduce the flight delays:

```
arrange(flights, desc(arr_delay))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     9       641             900         1301    1242
## 2  2013     6    15      1432            1935         1137    1607
## 3  2013     1    10      1121            1635         1126    1239
## 4  2013     9    20      1139            1845         1014    1457
## 5  2013     7    22       845            1600         1005    1044
## 6  2013     4    10      1100            1900          960    1342
## 7  2013     3    17      2321             810          911     135
## 8  2013     7    22      2257             759          898     121
## 9  2013    12     5       756            1700          896    1058
## 10 2013     5     3      1133            2055          878    1250
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

Deformation(calculation) function —-mutate

In the dplyr package, we can use the **mutate()** function to generate new variables directly from existing data, which is especially useful when using related classes and clustering algorithms.

```
mutate(flights, gain = arr_delay - dep_delay, speed = distance/air_time*60)
```

```
## # A tibble: 336,776 x 21
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1       517             515          2     830
## 2  2013     1     1       533             529          4     850
## 3  2013     1     1       542             540          2     923
## 4  2013     1     1       544             545         -1    1004
## 5  2013     1     1       554             600         -6     812
## 6  2013     1     1       554             558         -4     740
## 7  2013     1     1       555             600         -5     913
## 8  2013     1     1       557             600         -3     709
## 9  2013     1     1       557             600         -3     838
## 10 2013     1     1       558             600         -2     753
## # ... with 336,766 more rows, and 14 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>, gain <dbl>, speed <dbl>
```

```
# source: local data frame[336,776 x 21]
```

The code above generates 2 new variables and rows of **gain** and **speed**. Gain is equal to the delay time of leaving time minus the time delay of arrival time, and speed is equal to the distance divided by time * 60. From the output, we can see that these 2 columns have been added to data.frame (the bottom Variable not shown).

Also, we could use **transform()** function to modify the existing function directly to form new rows(variables.) But if you want to keep the new formed rows(variables) only, you can use **transmute()** function.

```
transmute(flights,
  gain = arr_delay - dep_delay,
  gain_per_hour = gain / (air_time / 60)
)
```

```
## # A tibble: 336,776 x 2
##   gain gain_per_hour
##   <dbl>         <dbl>
## 1     9      2.378855
## 2    16      4.229075
## 3    31     11.625000
## 4   -17     -5.573770
## 5   -19     -9.827586
## 6    16      6.400000
## 7    24      9.113924
## 8   -11    -12.452830
## 9    -5      -2.142857
## 10   10      4.347826
## # ... with 336,766 more rows
```

Summary function —-summarize

In dplyr package, we use the **summarize()** function to compile the data. The following code indicates that the average departure time delay is averaged, with the **na.rm** saying that all rows with missing data are removed.

```
summarise(flights,
  delay = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 1 x 1
##   delay
##   <dbl>
## 1 12.63907
```

```
# source: local data frame [1x1]
```

Besides, we can also use `sample_n()` and `sample_frac()` function to choose data randomly and calculate the exact part of data we want. This is very important, because the data gathering function, we can easily found the target data in the huge mountain. **### grouping function--group_by**

```
by_tailnum <- group_by(flights, tailnum)
delay <- summarise(by_tailnum,
  count = n(),
  dist = mean(distance, na.rm = TRUE),
  delay = mean(arr_delay, na.rm = TRUE))
delay <- filter(delay, count > 20, dist < 2000)
delay
```

```
## # A tibble: 2,962 x 4
##   tailnum count    dist    delay
##   <chr> <int>   <dbl>   <dbl>
## 1 N0EGMQ   371 676.1887  9.9829545
## 2 N10156   153 757.9477 12.7172414
## 3 N102UW    48 535.8750  2.9375000
## 4 N103US    46 535.1957 -6.9347826
## 5 N104UW    47 535.2553  1.8043478
## 6 N10575   289 519.7024 20.6914498
## 7 N105UW    45 524.8444 -0.2666667
## 8 N107US    41 528.7073 -5.7317073
## 9 N108UW    60 534.5000 -1.2500000
## 10 N109UW   48 535.8750 -2.5208333
## # ... with 2,952 more rows
```

Multi-step operation concatenations—`%>%`

The pipeline operator in R is `%>%`, and this symbol can link a series of action functions. The pipe operator `%>%` will connect the different code instructions. The `%>%` symbol will pass the output on the left to the right as the first parameter of the right function. It is useful to use the pipe operator when performing a continuous operations on a data set, which allows you to record the output of each operation not gradually.

```
flights2 <- mutate(flights, speed = distance / (air_time / 60))
speed <- select(flights2, tailnum, speed)
speed %>%
  group_by(tailnum) %>%
  summarise(count = n(), avg_speed = mean(speed, na.rm = TRUE)) %>%
  arrange(desc(avg_speed))
```

```
## # A tibble: 4,044 x 3
##   tailnum count avg_speed
##   <chr> <int>   <dbl>
## 1 N228UA     1 500.8163
## 2 N315AS     1 498.6851
## 3 N654UA     1 498.5821
## 4 N819AW     1 490.3448
## 5 N382HA    26 485.6026
## 6 N388HA    36 484.3891
## 7 N391HA    21 484.0645
## 8 N777UA     1 483.3645
## 9 N385HA    28 482.8947
## 10 N392HA    13 482.2468
## # ... with 4,034 more rows
```

Simple random sample function—`sample_n`, `sample_frac`

We can use `sample_n()` function and `sample_frac()` to take a random sample of rows: use `sample_n()` for a fixed number, and `sample_frac()` for a fixed fraction.

```
sample_n(flights, 10)
```

```
## # A tibble: 10 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013    12    14     1906           1620        166    2224
## 2  2013     2    18      900           905         -5    1218
## 3  2013    10    20     1329           1330         -1    1456
## 4  2013     5    24      927           932         -5    1227
## 5  2013    11     5      936           940         -4    1108
## 6  2013     2    28     1653           1655         -2    2005
## 7  2013    11    22     1627           1543         44    1851
## 8  2013     3     3     1357           1354          3    1633
## 9  2013    10    13     1619           1620         -1    1927
## 10 2013     1    21     2352           2359         -7     433
## # ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

```
sample_frac(flights,0.01)
```

```
## # A tibble: 3,368 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     9    20     1754           1800         -6    1850
## 2  2013    12    26     1219           1220         -1    1406
## 3  2013    11     4     1541           1547         -6    1820
## 4  2013    12    19     1857           1900         -3    2154
## 5  2013     2     5      754           753          1     NA
## 6  2013     4    23     1022           940         42    1315
## 7  2013    10     8     1449           1452         -3    1744
## 8  2013    11    12      703           705         -2     854
## 9  2013     3     9      918           900         18    1211
## 10 2013     1    29     1249           1300        -11    1345
## # ... with 3,358 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

use **replace = TRUE** to perform a bootstrap sample. we can also weight the sample with the **weight** argument.

** Using dplyr to analysis the relationship between flight distance and delay time**

Flight delays and cancellation are the case for everyone who has the needs of travel. Now let using *nycflights13* package to do the statistical analysis of flight data.we briefly studies the correlation between flight distance and delay time.

```
disDelay <- function(){
  myFlights <- select(flights,
    year,month,day,
    dep_delay,arr_delay,
    distance,dest)

  myFlights <- select(flights,
    year,month,day,
    dep_delay,arr_delay,
    distance,dest)

  myFlights

  # list the renaming names.
  myFlights <- rename(myFlights,destination = dest)
  # delete the missing data
  myFlights <- filter(myFlights,
    !is.na(dep_delay),
    !is.na(arr_delay))

  # data arrangment
  arrange(myFlights,dep_delay)
  arrange(myFlights,desc(dep_delay))
  # data calculation : the relationship between flight distance and delay time
  be_dest<- group_by(myFlights,be_dest)# data analysis\
  delay <- summarise(be_dest,      # Statistical calculation after grouping data
    count = n(),      # number of flights
    dist = mean(distance,na.rm = TRUE),
    delay = mean(arr_delay, na.rm = TRUE)
  )

  # remove distrubting data
  delay <- filter(delay,count > 20)
  return(delay)}
```

now we want to display data with visualized picture

```
library("ggplot2")
```



```
## Warning: package 'ggplot2' was built under R version 3.4.2
```

```
delay
```

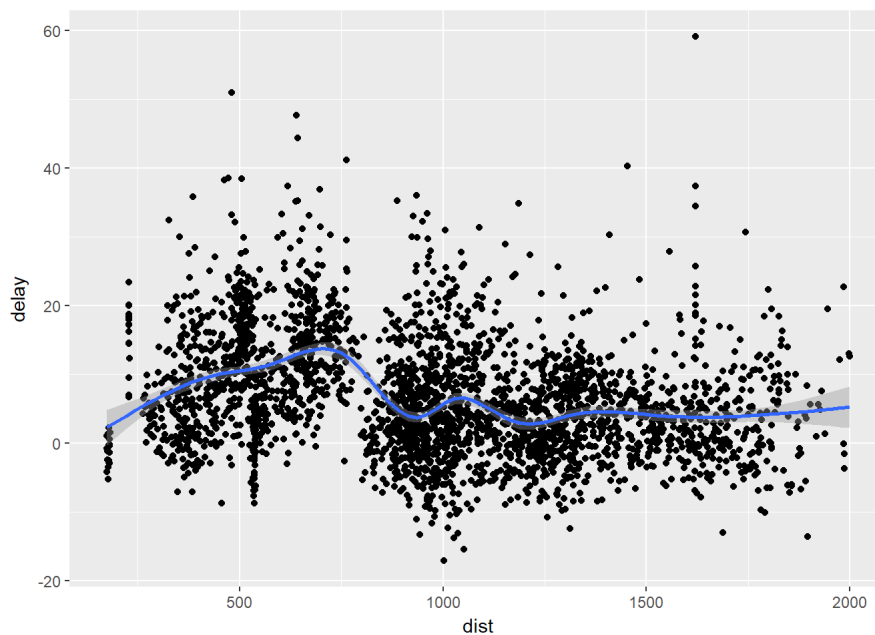
```
## # A tibble: 2,962 x 4
##   tailnum count    dist    delay
##   <chr> <int> <dbl> <dbl>
## 1 N0EGMQ   371 676.1887  9.9829545
## 2 N10156   153 757.9477 12.7172414
## 3 N102UW    48 535.8750  2.9375000
## 4 N103US    46 535.1957 -6.9347826
## 5 N104UW    47 535.2553  1.8043478
## 6 N10575   289 519.7024 20.6914498
## 7 N105UW    45 524.8444 -0.2666667
## 8 N107US    41 528.7073 -5.7317073
## 9 N108UW    60 534.5000 -1.2500000
## 10 N109UW   48 535.8750 -2.5208333
## # ... with 2,952 more rows
```

```
ggplot(data = delay) +
  geom_point(mapping = aes(x = dist, y = delay)) +
  geom_smooth(mapping = aes(x = dist, y = delay))
```

```
## `geom_smooth()` using method = 'gam'
```

```
## Warning: Removed 1 rows containing non-finite values (stat_smooth).
```

```
## Warning: Removed 1 rows containing missing values (geom_point).
```



conclusion

From the picture above, we can conclude :

1. Within 2500 miles, there is certain relationship between the aircraft distance and flight delays. There is basically no correlation between 2500 miles or more distance and flight delay times.
2. Flight delay times are basically within 20 minutes.

Further tools in dplyr that we did not cover in Class

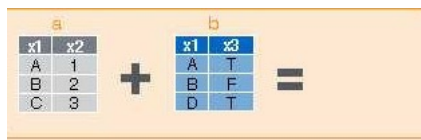
join function

When analysis data, we want to get the most efficient way. In **base** function, we could use **merge()** function to combine two data frames together. But when our population size are large, using **merge()**, it often takes ten or more minutes to get the output. Using **join()** in **dplyr** is our best choice now.

Join, like merge, is designed for the types of problems where you would use a [sql](#) join.

There are 6 types of join in dplyr. * **inner_join()** * **left_join()** * **right_join()** * **semi_join()** * **anti_join()** * **full_join()**

Suppose we have 2 tables A and B



Image

1. `inner_join(a,b,by = "x1")` -- Merge data, keep all records, all rows.
2. `left_join(a,b,by = "x1")` left_join(a,b,by = "x1")
3. `right_join(a,b,by = "x1")` -- Adding a matching data set B record to data set A
4. `semi_join(a,b,by = "x1")` -- Merge data, keep all records, all rows.
5. `anti_join(a,b,by = "x1")` -- Data set A does not match the data set B
6. `full_join(a,b,by = "x1")` -- Like merge() function.

Use 2 simple data frames to demonstrate how joins work.

```
monitors <- read.table(header=TRUE, text='
  monitorid      lat      long
    1  42.467573  -87.810047
    2  42.049148  -88.273029
    3  39.110539  -90.324080
  ')

pollutants <- read.table(header=TRUE, text='
  pollutant  duration  monitorid
    ozone      1h           1
    so2         1h           1
    ozone      8h           2
    no2         1h           4
  ')
```

```
library(dplyr)
```

```
inner_join(pollutants, monitors, by = "monitorid")
```

```
## pollutant duration monitorid      lat      long
## 1    ozone      1h           1 42.46757 -87.81005
## 2    so2        1h           1 42.46757 -87.81005
## 3    ozone      8h           2 42.04915 -88.27303
```

```
left_join(pollutants, monitors, by = "monitorid")
```

```
## pollutant duration monitorid      lat      long
## 1    ozone      1h           1 42.46757 -87.81005
## 2    so2        1h           1 42.46757 -87.81005
## 3    ozone      8h           2 42.04915 -88.27303
## 4    no2        1h           4      NA      NA
```

```
full_join(pollutants, monitors, by = "monitorid")
```

```
## pollutant duration monitorid      lat      long
## 1    ozone      1h           1 42.46757 -87.81005
## 2    so2        1h           1 42.46757 -87.81005
## 3    ozone      8h           2 42.04915 -88.27303
## 4    no2        1h           4      NA      NA
## 5    <NA>       <NA>           3 39.11054 -90.32408
```

```
semi_join(pollutants, monitors, by = "monitorid")
```

```
## pollutant duration monitorid
## 1    ozone      1h           1
## 2    so2        1h           1
## 3    ozone      8h           2
```

```
anti_join(monitors, pollutants, by = "monitorid")
```

```
## monitorid      lat      long
## 1         3 39.11054 -90.32408
```

Summary

We have learned most of the function of **dplyr** package in class, and we study the **join()** function today. The example I give is basic, we need to understand how convenient the **dplyr** package is, the more we operate on the data, the more advantages of **dplyr** will be gradually reflected.



Reference 1.<https://rpubs.com/NateByers/Merging>
2.<https://cran.r-project.org/web/packages/dplyr/vignettes/dplyr.html>
3.<https://github.com/andrewpbray/working-with-data-in-r/blob/master/working-with-data-in-r.Rmd>
4.<http://www.cnblogs.com/big-face/p/4863001.html>
5.<http://blog.csdn.net/wlt9037/article/details/74420886>
6.<http://www.jianshu.com/p/b2abad66cb01>
7.<https://wenku.baidu.com/view/53dd9590770bf78a64295460.html>
8.https://www.w3schools.com/sql/sql_ref_sqlserver.asp
9.<https://cran.r-project.org/web/packages/nycflights13/nycflights13.pdf>